

MIDS-W261-2016-HWK-Assignment4-Team4

February 27, 2016

=====
DATSCIW261 ASSIGNMENT #4
MIDS UC Berkeley, Machine Learning at Scale
DATSCIW261 ASSIGNMENT #4
Version 2016-01-27 (FINAL)

Hetal Chandaria, Patrick Ng, Marjorie Sayer W261 - 2 , ASSIGNMENT #4

Submission Date : Feb 12, 2016

Group : 4

HW 4.0. What is MrJob? How is it different to Hadoop MapReduce?

What are the `mapper_init`, `mapper_final`(), `combiner_final`(), `reducer_final`() methods? When are they called?

Answer

1. MrJob is a python package that helps us write and run Hadoop streaming jobs. It assists us in submitting job to Hadoop job tracker and in running each individual step under Hadoop streaming.

MapReduce is a framework processing parallelizable problems across huge data sets, using a large number of computers (nodes); cluster or grid. User specify a map function that processes a key/value pair to generate intermediate key/value pairs and a reduce function that merges all values associated with the same intermediate key.

Hadoop MapReduce is an implementation of Mapreduce programming framework. This API is implemented in Java. The Hadoop streaming library internally uses this. MrJob supports both local , hadoop and AWS EMR modes. For hadoop mode it internally uses hadoop streaming.

2.

`mapper_init()` : is used to define an action to be run before the mapper processes any input. `mapper_init()` is called first if there is function defined.

`mapper_final()`: is used to define an action to run after the mapper reaches the end of input. `mapper_final` is called after the mapper function has finished processing but before combiner or reducer is called.

`combiner_final()`: is used to define an action to run after the combiner reaches the end of input. `combiner_final` is called after the combiner function has finished processing but before the reducer is called.

`reducer_final()`: is used to define an action to run after the reducer reaches the end of input. `combiner_final` is called after the reducer function is finished.

HW 4.1 What is serialization in the context of MrJob or Hadoop?

When it used in these frameworks?

What is the default serialization mode for input and outputs for MrJob?

Answer What is serialization in the context of MrJob or Hadoop?

Serialization (and deserialization) implies being able to take an object, convert it into bytes and then be able to reconstruct the object back from those bytes. In the context of Hadoop, the mapreduce APIs act upon keys and values. For primitive types serialization would not be a major issue to overcome. However, for complex types, the MR framework needs to know how to take the raw input and convert it into the keys&values for the mappers, then take the keys and values generated by the mapper, serialize them as needed for the shuffle layer to make them available to the reducer and finally write the key-value output generated by the reducer to HDFS/disk in the format needed. This requires the mapreduce framework to understand the serialization format for these keys and values for the framework to function. In the context of MRJob, it uses json for the most part. This can be overridden via use of mrjob custom protocols.

When it used in these frameworks?

- a. De-serialization: Converting raw input to keys and values for the mapper.
- b. Serialize and de-serialize: Write mapper output to disk. Convert bytes received from shuffle layer to keys and values for reducer input.
- c. Serialization: Write reducer output to disk/HDFS.

What is the default serialization mode for input and outputs for MrJob?

The default input protocol for serialization is RawValueProtocol. The default output and internal protocol are both JSONProtocol.

HW 4.2: Recall the Microsoft logfiles data from the async lecture. The logfiles are described are located at:

<https://kdd.ics.uci.edu/databases/msweb/msweb.html>

<http://archive.ics.uci.edu/ml/machine-learning-databases/anonymouse/>

This dataset records which areas (Vroots) of www.microsoft.com each user visited in a one-week timeframe in February 1998.

Here, you must preprocess the data on a single node (i.e., not on a cluster of nodes) from the format:

```
C,"10001",10001 #Visitor id 10001
V,1000,1 #Visit by Visitor 10001 to page id 1000
V,1001,1 #Visit by Visitor 10001 to page id 1001
V,1002,1 #Visit by Visitor 10001 to page id 1002
C,"10002",10002 #Visitor id 10001
V
```

Note: #denotes comments to the format:

```
V,1000,1,C, 10001
V,1001,1,C, 10001
V,1002,1,C, 10001
```

Write the python code to accomplish this.

```
In [3]: %%writefile preprocess_log.py
        #!/usr/bin/python
        """ This program process the log file and outputs 2 separate files.
            One for the log information and other is page URLS's.
        """
        import sys
        import re

        inputfile = sys.argv[1]
        outputfile = open('processed_anonymous-msweb.data', 'a')
        outputfile2 = open('processed_urls.data', 'a')

        for line in open(inputfile):
```

```

v = line.split(',')
# If the line is customer information line, save it to customer_info
if(v[0]=='C'):
    customer_info = v
# If the line visit information line, concatenate the line with customer_info
elif(v[0]=='V'):
    outputfile.write( line.strip()+','+customer_info[0]+','+customer_info[1].strip('"')+'\n')
# Directly output other lines
elif(v[0]=='A'):
    outputfile2.write(v[1]+','+ re.sub(r'\'+',' ',v[4]))
else:
    continue

```

Writing preprocess_log.py

```

In [4]: !chmod a+x preprocess_log.py
        ! rm processed_anonymous-msweb.data
        ! rm processed_urls.data
        !./preprocess_log.py anonymous-msweb.data
        !wc -l processed_anonymous-msweb.data
        ! head processed_anonymous-msweb.data
        !wc -l processed_urls.data
        ! head processed_urls.data

```

rm: processed_anonymous-msweb.data: No such file or directory

rm: processed_urls.data: No such file or directory

98654 processed_anonymous-msweb.data

V,1000,1,C,10001

V,1001,1,C,10001

V,1002,1,C,10001

V,1001,1,C,10002

V,1003,1,C,10002

V,1001,1,C,10003

V,1003,1,C,10003

V,1004,1,C,10003

V,1005,1,C,10004

V,1006,1,C,10005

294 processed_urls.data

1287,/autoroute

1288,/library

1289,/masterchef

1297,/centroam

1215,/developer

1279,/msgolf

1239,/msconsult

1282,/home

1251,/referencesupport

1121,/magazine

HW 4.3: Find the 5 most frequently visited pages using MrJob from the output of 4.2 (i.e., transformed log file).

```

In [1]: %load_ext autoreload
        %autoreload 2

```

```
In [5]: %%writefile hw4_3.py
```

```
'''
Input format:
....
V,1000,1,C,10001
V,1001,1,C,10001
V,1002,1,C,10001
V,1001,1,C,10002
V,1003,1,C,10002
V,1001,1,C,10003
V,1003,1,C,10003
V,1004,1,C,10003
....
'''

from mrjob.job import MRJob, MRStep
import csv
import sys

def csv_readline(line):
    """Given a sting CSV line, return a list of strings."""
    return csv.reader([line]).next()

class PageVisitHW4_3(MRJob):

    def mapper_get_visit_count(self, line_no, line):
        """Extracts the page id and visit count"""
        cell = csv_readline(line)
        yield cell[1], 1

    def reducer_get_visit_count(self, pageId, visit_counts):
        """Sumarizes the visit counts by adding them together."""
        total = sum(i for i in visit_counts)
        yield pageId, total

    def reducer_find_top5_pages_init(self):
        self.printed = 0

    def reducer_find_top5_pages(self, pageId, visit_counts):
        """Print the top 5 pageId's and the counts"""
        if self.printed < 5:
            yield pageId, visit_counts.next()
            self.printed += 1

    def steps(self):
        return [
            MRStep(mapper=self.mapper_get_visit_count,
                    combiner=self.reducer_get_visit_count,
                    reducer=self.reducer_get_visit_count),
            MRStep(reducer_init=self.reducer_find_top5_pages_init,
                    reducer=self.reducer_find_top5_pages,
                    jobconf={
                        "stream.num.map.output.key.fields": "2",
                        "mapreduce.job.output.key.comparator.class":
```

```

        "org.apache.hadoop.mapred.lib.KeyFieldBasedComparator",
        "mapreduce.partition.keycomparator.options": "-k2,2nr"
    }])

```

```

if __name__ == '__main__':
    PageVisitHW4_3.run()

```

Overwriting hw4_3.py

```

In [12]: from hw4_3 import PageVisitHW4_3
        mr_job = PageVisitHW4_3(args=['processed_anonymous-msweb.data', '-r', 'hadoop', '--strict-prot
        with mr_job.make_runner() as runner:
            runner.run()
            # stream_output: get access of the output

            print "5 most frequently visited pages:"

            for line in runner.stream_output():
                print mr_job.parse_output_line(line)

```

```

5 most frequently visited pages:
('1008', 10836)

```

```

ERROR:mrjob.fs.hadoop:STDERR: 16/02/11 11:57:19 WARN util.NativeCodeLoader: Unable to load native-hadoop

```

```

('1034', 9383)
('1004', 8463)
('1018', 5330)
('1017', 5108)

```

HW 4.4: Find the most frequent visitor of each page using MrJob and the output of 4.2 (i.e., transformed log file). In this output please include the webpage URL, webpageID and Visitor ID.

```

In [8]: %%writefile hw4_4.py

```

```

from mrjob.job import MRJob, MRStep
import mrjob
import csv
import sys

def csv_readline(line):
    """Given a sting CSV line, return a list of strings."""
    return csv.reader([line]).next()

class PageVisitHW4_4(MRJob):

    BASE_URL = "http://www.microsoft.com"

    # Have to use Raw Protocol in order to get sorting (using 3rd key) to work.
    INTERNAL_PROTOCOL = mrjob.protocol.RawProtocol
    OUTPUT_PROTOCOL = mrjob.protocol.RawProtocol

    def mapper1(self, line_no, line):
        cell = csv_readline(line)

```

```

        # page,visitor \t count
        yield ",".join([cell[1], cell[4]]), "1"

    def reducer1(self, key, values):
        """Sum up the visit count per (page, visitor) pair."""
        total = sum([int(v) for v in values])
        fields = key.split(",")
        # page \t visitor \t total
        yield fields[0], "\t".join([fields[1], str(total)])

    def reducer2_init(self):
        # Build the dictionary of pageId:url
        self.urls = {}
        with open("processed_urls.data", "r") as f:
            for fields in csv.reader(f):
                self.urls[fields[0]] = fields[1]

    def reducer2(self, key, values):
        # url \t pageId \t visitor \t total
        yield self.BASE_URL + self.urls[key] + "\t" + key, values.next()

    def steps(self):
        return [
            MRStep(mapper=self.mapper1,
                    reducer=self.reducer1,
                    ),
            MRStep(reducer_init=self.reducer2_init,
                    reducer=self.reducer2,
                    jobconf={
                        "stream.num.map.output.key.fields":"3",
                        "mapreduce.job.output.key.comparator.class":
                            "org.apache.hadoop.mapred.lib.KeyFieldBasedComparator",
                        "mapreduce.partition.keycomparator.options":"-k3,3nr"
                    })]

if __name__ == '__main__':
    PageVisitHW4_4.run()

```

Overwriting hw4_4.py

```

In [13]: !python ./hw4_4.py \
        -r hadoop \
        --file processed_urls.data \
        --strict-protocols \
        processed_anonymous-msweb.data > most_visitors.out

```

```

no configs found; falling back on auto-configuration
no configs found; falling back on auto-configuration
creating tmp directory /var/folders/dm/nsw7wjf91f1c74hgl17ldw040000gn/T/hw4_4.patrickng.20160211.035741.
writing wrapper script to /var/folders/dm/nsw7wjf91f1c74hgl17ldw040000gn/T/hw4_4.patrickng.20160211.035741.
Using Hadoop version 2.7.1
Copying local files into hdfs:///user/patrickng/tmp/mrjob/hw4_4.patrickng.20160211.035741.553883/files/
HADOOP: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
HADOOP: packageJobJar: [/var/folders/dm/nsw7wjf91f1c74hgl17ldw040000gn/T/hadoop-unjar9199386722825397230

```

```

Counters from step 1:
(no counters found)
HADOOP: Unable to load native-hadoop library for your platform... using builtin-java classes where appl
HADOOP: packageJobJar: [/var/folders/dm/nsw7wjf91fc74hgl17ldw040000gn/T/hadoop-unjar457463116393516219
Counters from step 2:
(no counters found)
Streaming final output from hdfs:///user/patrickng/tmp/mrjob/hw4.4.patrickng.20160211.035741.553883/outp
STDERR: 16/02/11 11:59:01 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your pla

removing tmp directory /var/folders/dm/nsw7wjf91fc74hgl17ldw040000gn/T/hw4.4.patrickng.20160211.035741.553883
deleting hdfs:///user/patrickng/tmp/mrjob/hw4.4.patrickng.20160211.035741.553883 from HDFS

```

```

In [14]: print "Webpage URL, Page Id, Customer Id, Visit Count"
!head most_visitors.out | sed s/"//g

```

Webpage URL	Page Id	Customer Id	Visit Count
http://www.microsoft.com/train_cert	1295	42616	1
http://www.microsoft.com/partner	1284	41108	1
http://www.microsoft.com/cinemania	1283	41033	1
http://www.microsoft.com/home	1282	41244	1
http://www.microsoft.com/intellimouse	1281	37099	1
http://www.microsoft.com/music	1280	41643	1
http://www.microsoft.com/msgolf	1279	31062	1
http://www.microsoft.com/hed	1278	41317	1
http://www.microsoft.com/stream	1277	30111	1
http://www.microsoft.com/vtestsupport	1276	40810	1

HW 4.5 Here you will use a different dataset consisting of word-frequency distributions for 1,000 Twitter users. These Twitter users use language in very different ways, and were classified by hand according to the criteria:

- 0: Human, where only basic human-human communication is observed.
- 1: Cyborg, where language is primarily borrowed from other sources (e.g., jobs listings, classifieds postings, advertisements, etc...).
- 2: Robot, where language is formulaically derived from unrelated sources (e.g., weather/seismology, police/fire event logs, etc...).
- 3: Spammer, where language is replicated to high multiplicity (e.g., celebrity obsessions, personal promotion, etc...)

Check out the preprints of our recent research, which spawned this dataset:

<http://arxiv.org/abs/1505.04342>

<http://arxiv.org/abs/1508.01843>

The main data lie in the accompanying file:

topUsers_Apr-Jul_2014_1000-words.txt

and are of the form:

USERID,CODE,TOTAL,WORD1.COUNT,WORD2.COUNT,... . .

where

USERID = unique user identifier CODE = 0/1/2/3 class code TOTAL = sum of the word counts

Using this data, you will implement a 1000-dimensional K-means algorithm in MrJob on the users by their 1000-dimensional word stripes/vectors using several centroid initializations and values of K.

Note that each “point” is a user as represented by 1000 words, and that word-frequency distributions are generally heavy-tailed power-laws (often called Zipf distributions), and are very rare in the larger class of discrete, random distributions. For each user you will have to normalize by its “TOTAL” column. Try several parameterizations and initializations:

- (A) K=4 uniform random centroid-distributions over the 1000 words

- (B) K=2 perturbation-centroids, randomly perturbed from the aggregated (user-wide) distribution
- (C) K=4 perturbation-centroids, randomly perturbed from the aggregated (user-wide) distribution
- (D) K=4 “trained” centroids, determined by the sums across the classes.

and iterate until a threshold (try 0.001) is reached. After convergence, print out a summary of the classes present in each cluster. In particular, report the composition as measured by the total portion of each class type (0-3) contained in each cluster, and discuss your findings and any differences in outcomes across parts A-D.

Note that you do not have to compute the aggregated distribution or the class-aggregated distributions, which are rows in the auxiliary file:

topUsers_Apr-Jul_2014_1000-words_summaries.txt

```
In [262]: %%writefile custom_func.py
# function to calculate purity of cluster and print it. Its a generic function that can be ca

def calc_purity(cluster_dist):
    #calculate purity and print class distribution
    print 'Cluster distribution'
    print '-'*100

    user_class = { 0:'Human', 1:'Cyborg', 2:'Robot', 3:'Spammer' }
    human = sum([cluster_dist[k].get('0',0) for k in cluster_dist.keys()])
    cyborg = sum([cluster_dist[k].get('1',0) for k in cluster_dist.keys()])
    robot = sum([cluster_dist[k].get('2',0) for k in cluster_dist.keys()])
    spammer = sum([cluster_dist[k].get('3',0) for k in cluster_dist.keys()])
    print "{0:>5} |{1:>15} |{2:>15} |{3:>15} |{4:>15}".format(
        "k", "Human"+' ':''+str(human), "Cyborg"+' ':''+str(cyborg), "Robot"+' ':''+str(robot), "Spammer"+' ':''+str(spammer))
    print '-'*100
    max_cl={}
    total = 0
    for cid, cvalue in cluster_dist.iteritems():
        total += sum(cvalue.values())
        print "{0:>5} |{1:>15} |{2:>15} |{3:>15} |{4:>15}".format(
            cid, cvalue.get('0',0) ,cvalue.get('1',0) ,cvalue.get('2',0) , cvalue.get('3',0))
        max_cl[cid]=max(cvalue.values())
    print '-'*100
    print 'purity : %3.3f' %(100*sum(max_cl.values())/total)
    print '-'*100
```

Overwriting custom_func.py

```
In [263]: %%writefile Kmeans.py
from numpy import argmin, array, random
from mrjob.job import MRJob
from mrjob.step import MRJobStep
from itertools import chain
import numpy as np

#Calculate find the nearest centroid for data point
def MinDist(datapoint, centroid_points):
    datapoint = array(datapoint)
    centroid_points = array(centroid_points)
    diff = datapoint - centroid_points
    diffsq = diff*diff
```



```

        # Get the nearest centroid for each instance
        minidx = argmin(list(diffsq.sum(axis = 1)))
        return minidx

#Check whether centroids converge
def stop_criterion(centroid_points_old, centroid_points_new,T):
#    return np.alltrue(abs(np.array(centroid_points_new) - np.array(centroid_points_old)) <=
        oldvalue = list(chain(*centroid_points_old))
        newvalue = list(chain(*centroid_points_new))
        Diff = [abs(x-y) for x, y in zip(oldvalue, newvalue)]
        Flag = True
        for i in Diff:
            if(i>T):
                Flag = False
                break
        return Flag

class MRKmeans(MRJob):
    centroid_points=[]
    def steps(self):
        return [
            MRJobStep(mapper_init=self.mapper_init,mapper=self.mapper,combiner=self.combiner,
                ]

    #load centroids info from file
    def mapper_init(self):
        self.centroid_points = [map(float,s.split('\n')[0].split(',') for s in open("Centroid

    #load data and output the nearest centroid index and data point
    def mapper(self, _, line):
        D = line.split(',')
        total = int(D[2])
        code=int(D[1])
        #normalise the input data
        data=map(float,D[3:])
        normalise_data=map(lambda x:((1.0 * x) / total),data)
        #sending the class composition along with the mapper output
        yield int(MinDist(normalise_data,self.centroid_points)),(list(normalise_data),1,{code

    #Combine sum of data points locally
    def combiner(self, idx, inputdata):
        code = {} #for class composition
        sum_features= None
        num = 0 #for count
        for features,n,codes in inputdata:
            features = array(features) #convert to array first
            if sum_features is None: #if looping through first, initialise array of zeros
                sum_features = np.zeros(features.size)
            sum_features += features #add features
            num += n #increment count
            #count codes
            for k,v in codes.iteritems(): #increment class composition
                code[k] = code.get(k,0)+ v

```

```

        yield idx,(list(sum_features),num,code)

#Aggregate sum for each cluster and then calculate the new centroids
def reducer(self, idx, inputdata):
    centroids = None
    code = {}
    num = 0
    for features,n,codes in inputdata:
        features = array(features)

        if centroids is None:
            centroids = np.zeros(features.size)
            centroids += features

        num += n #increment predicted cluster count

    #count codes
    for k,v in codes.iteritems(): #increment class composition
        code[k] = code.get(k,0)+v

    centroids_new = centroids / num #compute new centroid
    yield idx, (list(centroids_new),code)

if __name__ == '__main__':
    MRKmeans.run()

```

Overwriting Kmeans.py

In [264]: `%%writefile run_kmeans.py`

```

from numpy import random
import numpy as np
from Kmeans import MRKmeans, stop_criterion
import sys
from custom_func import calc_purity
mr_job = MRKmeans(args=['topUsers_Apr-Jul_2014_1000-words.txt'])

random.seed(0)
#number of features
n= 1000

#get centroid type and number of clusters from user
if len(sys.argv) >2: k = int(sys.argv[2])
cen_type = sys.argv[1]

#Geneate initial centroids
centroid_points = []

#based on the centroid type generate centroids
if(cen_type=='Uniform'):
    rand_int = random.uniform(size=[k,n])
    total = np.sum(rand_int,axis=1)
    centroid_points = (rand_int.T/total).T
    with open('Centroids.txt', 'w+') as f:

```

```

        f.writelines(','.join(str(j) for j in i) + '\n' for i in centroid_points)
    f.close()

elif(cen_type=='Perturbation'):
    data = [s.split('\n')[0].split(',') for s in
            open("topUsers_Apr-Jul_2014_1000-words_summaries.txt").readlines()][1]
    #get the total count of words
    total = int(data[2])
    feature = map(lambda x:((1.0 * float(x)) / total),data[3:]) #normalise
    pertubation = feature + random.sample(size=(k,n)) #generate random sample and add it to f
    sum_per = np.sum(pertubation,axis=1) # calculate the sum to be used for normalization
    centroid_points = (pertubation.T/sum_per).T
    with open('Centroids.txt', 'w+') as f:
        f.writelines(','.join(str(j) for j in i) + '\n' for i in centroid_points)
    f.close()

else:
    data = [s.split('\n')[0].split(',') for s in
            open("topUsers_Apr-Jul_2014_1000-words_summaries.txt").readlines()][2:]
    for cluster in data:
        total = 0
        total = int(cluster[2])#get the total count of words
        feature = map(lambda x:((1.0 * float(x)) / total),cluster[3:]) #normalise
        centroid_points.append(feature)
    with open('Centroids.txt', 'w+') as f:
        f.writelines(','.join(str(j) for j in i) + '\n' for i in centroid_points)
    f.close()

print 'Centroid Type: %s' %cen_type

# Update centroids iteratively
i = 0
while(1):
    # save previous centroids to check convergency
    centroid_points_old = centroid_points[:]
    print "iteration"+str(i)+": "
    with mr_job.make_runner() as runner:
        centroid_points = []
        cluster_dist = {}
        runner.run()
        # stream_output: get access of the output
        for line in runner.stream_output():
            key,value = mr_job.parse_output_line(line)
            centroid, codes = value
            centroid_points.append(centroid)
            cluster_dist[key]=codes
    i = i + 1

    #check if we have convergence
    if(stop_criterion(centroid_points_old,centroid_points,0.001)):
        break

    #write new centroids back to file
    with open('Centroids.txt', 'w') as f:

```

```

        for centroid in centroid_points:
            f.writelines(','.join(map(str, centroid)) + '\n')
        f.close()

```

```

    calc_purity(cluster_dist)

```

Overwriting run_kmeans.py

In [248]: !python run_kmeans.py Uniform 4

Centroid Type: Uniform

iteration0:

iteration1:

iteration2:

iteration3:

iteration4:

iteration5:

Cluster distribution

k	Human:752	Cyborg:91	Robot:54	Spammer:103
0	0	0	11	0
1	0	51	0	0
2	1	37	38	4
3	751	3	5	99

purity : 85.100

In [249]: !python run_kmeans.py Perturbation 2

Centroid Type: Perturbation

iteration0:

iteration1:

iteration2:

iteration3:

Cluster distribution

k	Human:752	Cyborg:91	Robot:54	Spammer:103
0	751	3	16	99
1	1	88	38	4

purity : 83.900

In [250]: !python run_kmeans.py Perturbation 4

Centroid Type: Perturbation

iteration0:

iteration1:

iteration2:

iteration3:

iteration4:

iteration5:

iteration6:

Cluster distribution

k	Human:752	Cyborg:91	Robot:54	Spammer:103
0	0	0	5	0
1	0	51	0	0
2	1	37	38	4
3	751	3	11	99

purity : 84.500

In [247]: !python run_kmeans.py Trained

Centroid Type: Trained

iteration0:

iteration1:

iteration2:

iteration3:

iteration4:

Cluster distribution

k	Human:752	Cyborg:91	Robot:54	Spammer:103
0	749	3	14	38
1	0	51	0	0
2	1	37	40	4
3	2	0	0	61

purity : 90.100

In []:

HW4.6 (OPTIONAL) Scaleable K-MEANS++ Read the following paper entitled “Scaleable K-MEANS++” located at:

<http://theory.stanford.edu/~sergei/papers/vldb12-kmpar.pdf>

In MrJob, implement K-MEANS|| and compare with a random initialization for the dataset above. Report on the number passes over the training data, and time required to run all clustering algorithms.

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In [166]:

In []: